Chapter 7.

Indexing









Department of Computer Science, The University of Hong Kong

In this chapter...

- Outcome 1. Information Modeling
 - Able to understand the modeling of real life information in a database system.
- Outcome 2. Query Languages
 - Able to understand and use the languages designed for data access.
- Outcome 3. System Design
 - Able to understand the design of an efficient and reliable database system.
- Outcome 4. Application Development
 - Able to implement a practical application on a real database.

We are going to learn...

- Basic concepts
- B⁺ -tree
- Hashing
- Index definition in SQL

Section 1

Basic

Concepts

Basic concepts

- Index is used to speed up access to desired data.
 - E.g., Author catalog in library, phone directory index, etc.
- Search key
 - An attribute or a set of attributes used to look up records in a file.
- Indices are typically much smaller than the original file.



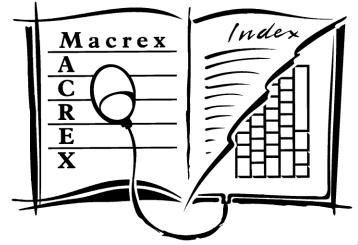
Primary v.s. secondary

- Primary index An index whose search key also defines the sequential order of the file.
 - E.g., Access staff records through staffID (primary search key).
- However, the data file can be sorted in one order only.
- How about accessing data with a different search key?
 - E.g., Access staff records through roomID (a secondary search key, need a secondary index!).

| staffID | roomID | faculty |
|---------|--------|---------|
| 10101 | 49 | C.S. |
| 12121 | 42 | Finance |
| 15151 | 35 | Music |
| 22222 | 10 | Physics |
| 32343 | 15 | History |
| 33456 | 18 | C.S. |
| 45565 | 20 | E.E.E. |
| 58583 | 3 | Biology |
| 76543 | 31 | Finance |
| 76766 | 5 | Finance |
| 83821 | 2 | C.S. |
| 98345 | 24 | C.S. |

2 classes of indices

- Ordered Indices Search keys are sorted in the index
 - Example: indexed-sequential file, B+-tree.
- Hash Indices Search keys are distributed over different buckets using a hash function.
 - Example: extendable hash-index.
- These indices differ in their speeds in answering different queries.



Index evaluation factors

Each indexing technique must be evaluated on the basis of these factors



- Access types The types of access that are supported efficiently (e.g., equality search or range search? Single attribute search or multi-attribute search?)
- Access time The time it takes to find a particular data item, or a set of items.
- Insertion / deletion time
- Space overhead

No one indexing technique is the best. Rather, each technique is best suited to particular database applications.



Section 2

B+-tree

Properties of B+-tree

- B+-tree index structure is one of the most widely used index structure in DBMS.
- All paths from root to leaf are of the same length (i.e., balanced)
- Can support efficient processing of the following queries (Assume that the B+-tree is built on attribute A of the relation R):

SELECT * FROM R WHERE R.A = 3

SELECT * FROM R WHERE R.A >= 3 AND R.A < 22

Why not binary search tree?



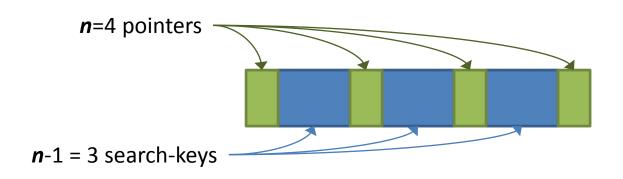
Balanced Binary Search tree minimizes the number of key comparisons for finding a search key. Why don't we use balanced binary search tree in Database?

- We need a tree which is
 - Node size = 1 block
 (A node can contain more than one keys)
 - Low in height
 - Balanced

the number of block retrieval in answering a query (i.e., number of tree nodes to be accessed) rather than the number of key comparisons.



A node in B+-tree

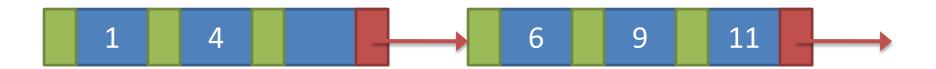


A node contains up to n-1 search-key values, and n pointers.



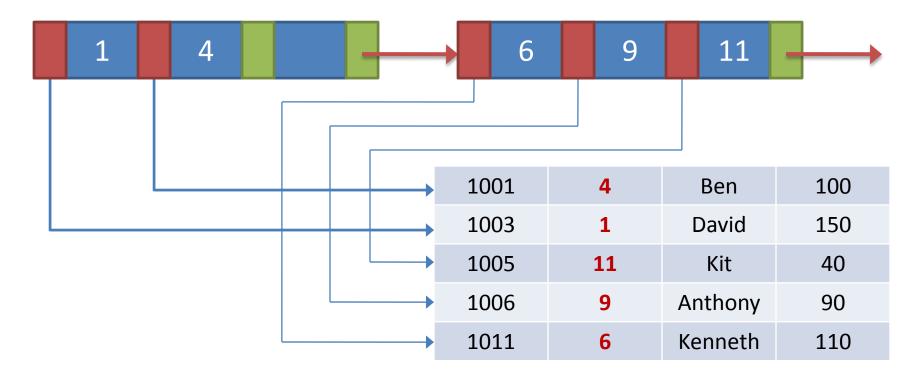
The search-key values within a node are kept in sorted order.

1. Leaf node



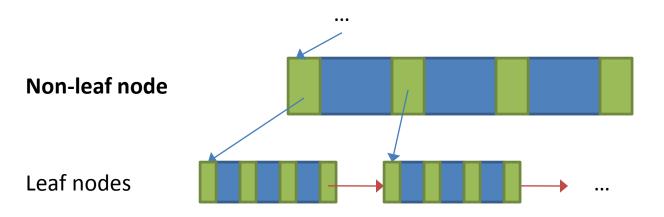
- A leaf node has at least $\lceil (n-1)/2 \rceil$ and at most (n-1) values, where n is the number of pointers.
 - E.g., with n = 4, a leaf node must contain at least 2 values, and at most 3 values.
- The last pointer is used to chain together the leaf nodes in search-key order.

1. Leaf node



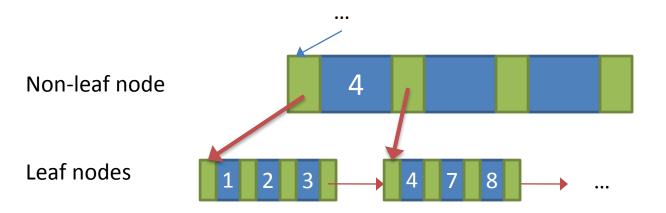
The pointer before a search-key value points to the record that contains the search-key (assume no duplicate values).

2. Non-leaf node



- Non-leaf nodes must hold at least $\lceil n/2 \rceil$, and at most n pointers.
 - E.g., with n = 4, a non-leaf node contains at least 2 pointers, and at most 4 pointers.

2. Non-leaf node



- The pointer on the left of a key K points to the part of the subtree that contains those key values less than K.
- The pointer on the right of a key K points to the part of the subtree that contains those key values larger than or equal to K.

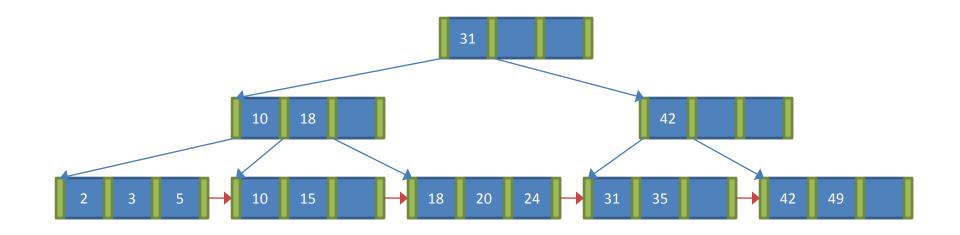
Example B+-tree

In the file, records are ordered according to the 1st attribute, we would like to **build a B⁺-tree index (secondary index)** to speed up the searching on the 2nd attribute.



| 10101 | 49 | C.S. |
|-------|----|---------|
| 12121 | 42 | Finance |
| 15151 | 35 | Music |
| 22222 | 10 | Physics |
| 32343 | 15 | History |
| 33456 | 18 | C.S. |
| 45565 | 20 | E.E.E. |
| 58583 | 3 | Biology |
| 76543 | 31 | Finance |
| 76766 | 5 | Finance |
| 83821 | 2 | C.S. |
| 98345 | 24 | C.S. |
| | | |

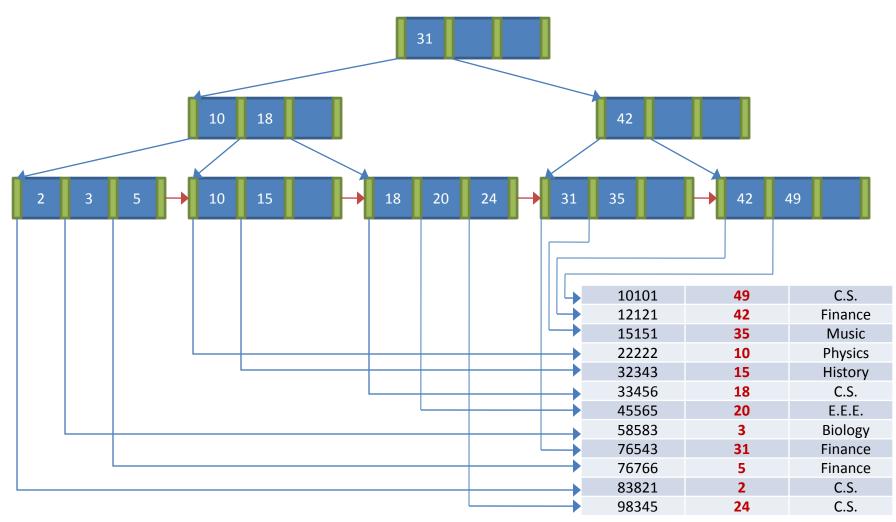
Example B+-tree



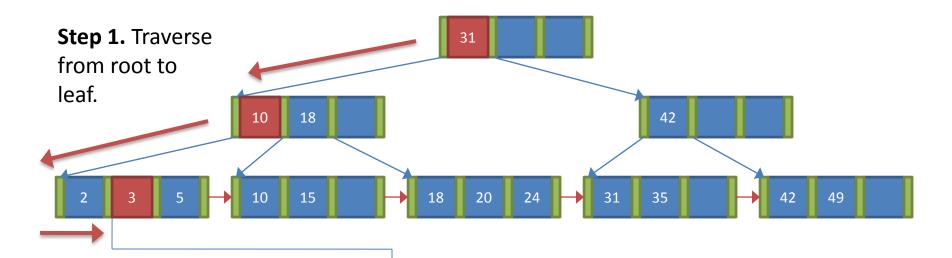
- With n = 4, a leaf node must contain at least 2 values, and at most 3 values.
- With n = 4, a non-leaf node must contain at least 2 pointers, and at most 4 pointers.

| 10101 | 49 | C.S. |
|-------|----|---------|
| 12121 | 42 | Finance |
| 15151 | 35 | Music |
| 22222 | 10 | Physics |
| 32343 | 15 | History |
| 33456 | 18 | C.S. |
| 45565 | 20 | E.E.E. |
| 58583 | 3 | Biology |
| 76543 | 31 | Finance |
| 76766 | 5 | Finance |
| 83821 | 2 | C.S. |
| 98345 | 24 | C.S. |
| | | |

Example B+-tree



Searching



Step 2. Search in the leaf node.

Point query

SELECT * FROM R WHERE R.B = 3

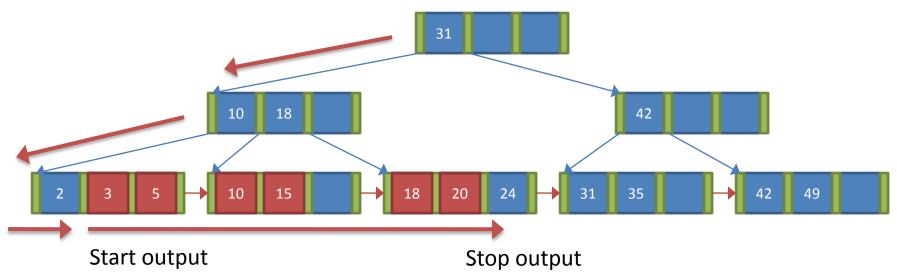
Step 3. Follow the pointer in the leaf node to retrieve the record.

| 10101 | 49 | C.S. |
|-------|----|---------|
| 12121 | 42 | Finance |
| 15151 | 35 | Music |
| 22222 | 10 | Physics |
| 32343 | 15 | History |
| 33456 | 18 | C.S. |
| 45565 | 20 | E.E.E. |
| 58583 | 3 | Biology |
| 76543 | 31 | Finance |
| 76766 | 5 | Finance |
| 83821 | 2 | C.S. |
| 98345 | 24 | C.S. |
| | | |



With this B+-tree, how many disk block accesses to answer this query?

Searching



Range query

SELECT * FROM R WHERE R.B >= 3 AND R.B < 22



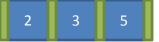
B+-tree can also handle **range search** very well. Search for the left border of the range and **traverse the leaf chain** until a record with search-key larger than the right border is encountered.

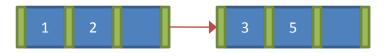
| 10101 | 49 | C.S. |
|-------|----|---------|
| 12121 | 42 | Finance |
| 15151 | 35 | Music |
| 22222 | 10 | Physics |
| 32343 | 15 | History |
| 33456 | 18 | C.S. |
| 45565 | 20 | E.E.E. |
| 58583 | 3 | Biology |
| 76543 | 31 | Finance |
| 76766 | 5 | Finance |
| 83821 | 2 | C.S. |
| 98345 | 24 | C.S. |
| | | |

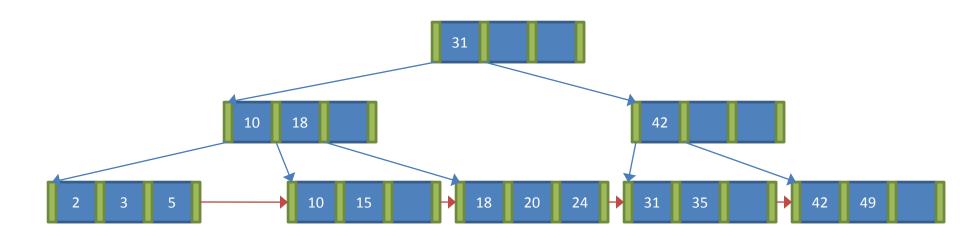
Insertion

- Assume no duplicate entries are inserted, insertion is simply searching + insert entry.
- If a leaf node is full, node splitting has to be performed.
 - Step 1. Create one more node and distribute the first $\lceil n/2 \rceil$ records to one node and the remaining to the other node.
 - Step 2. Parent nodes (non-leaf nodes) have to be updated accordingly.

Insert key "1"

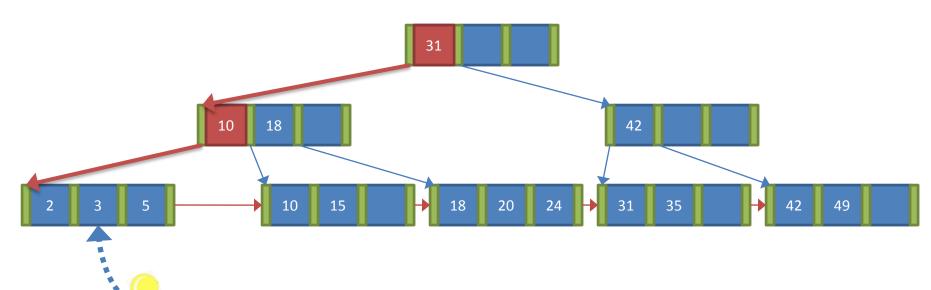






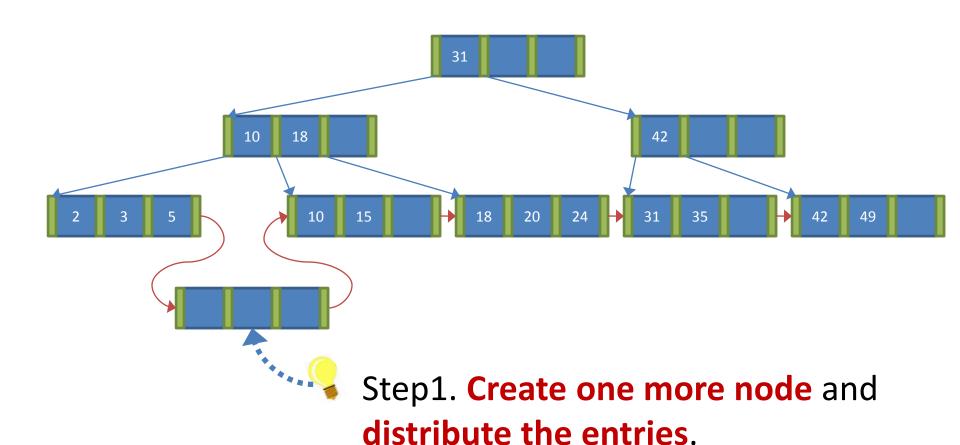


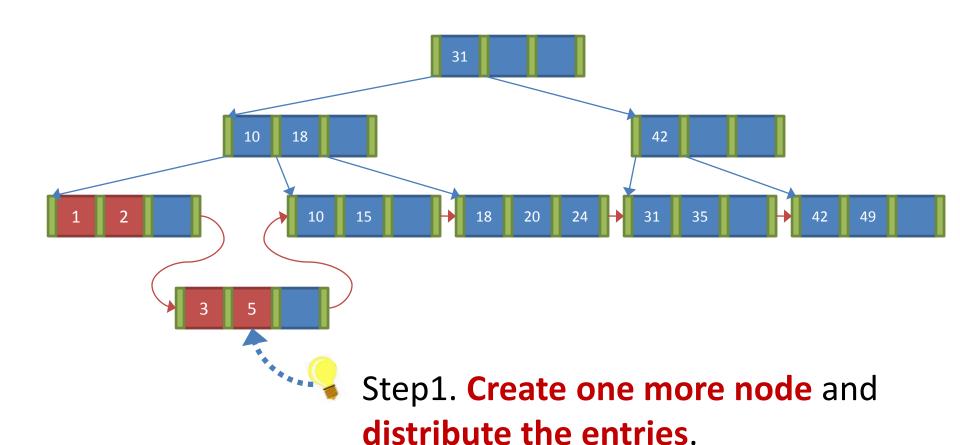
Let's learn how **node splitting** is implemented on leaf node by considering inserting key "1" in the above B+-tree.

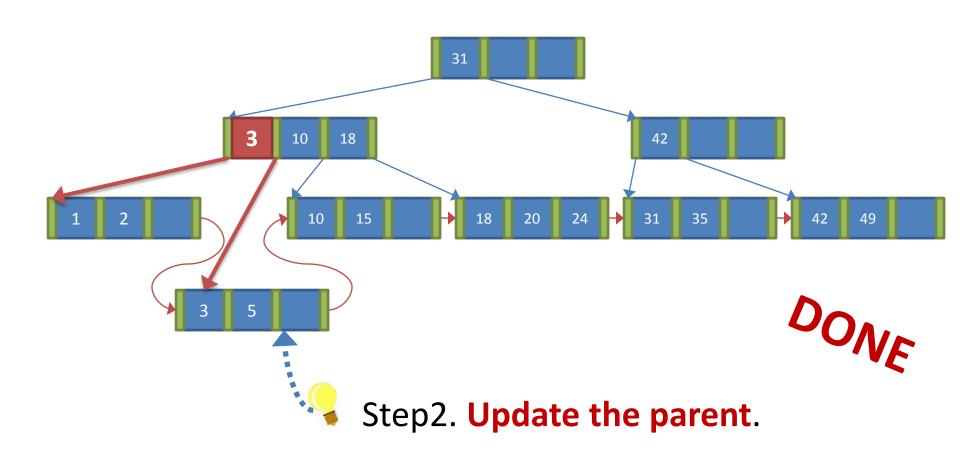


We first search for the leaf node that the key "1" should be inserted.

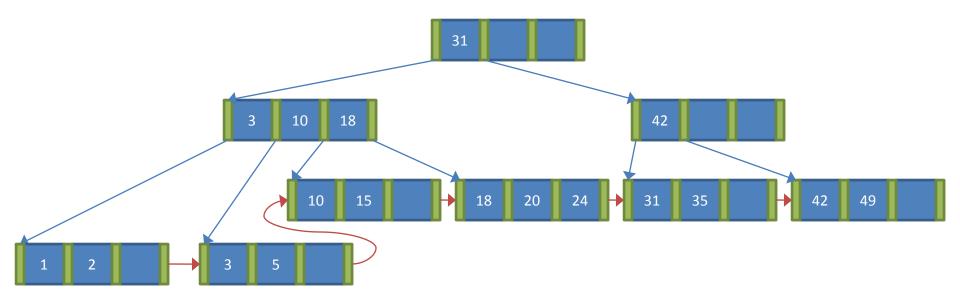
Since this node is **full**, inserting **"1"** requires **SPLITTING** this leaf node.





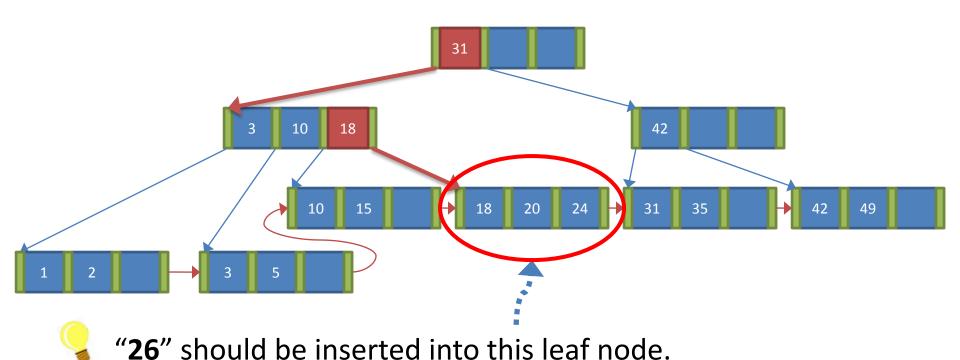


Splitting of a non-leaf node is a little different from splitting of a leaf node.

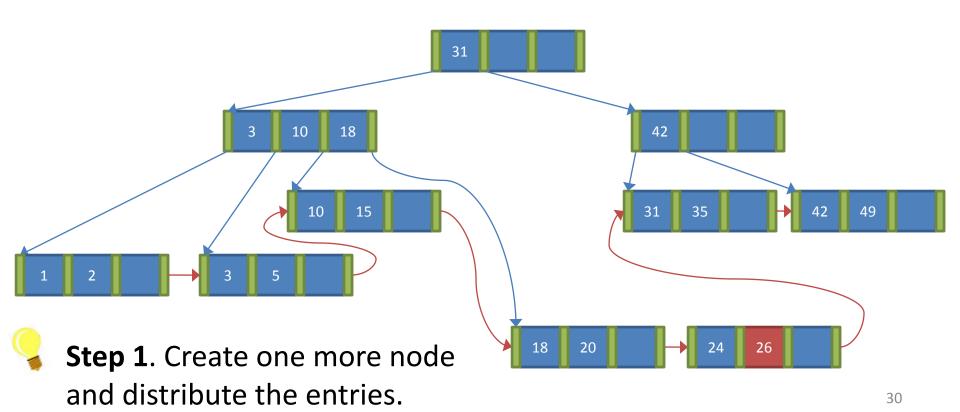




Let's learn how node splitting is implemented on non-leaf node by considering inserting key "26" in the above B+-tree.

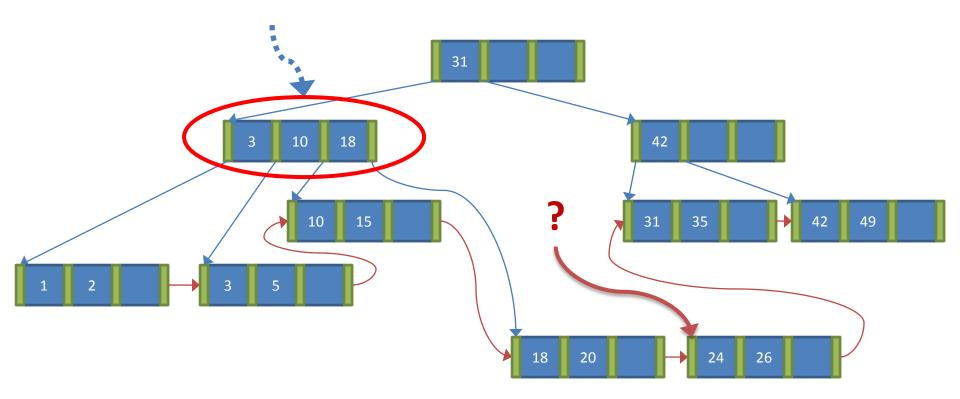


Since this node is **full**, node **SPLITTING** need to be performed.





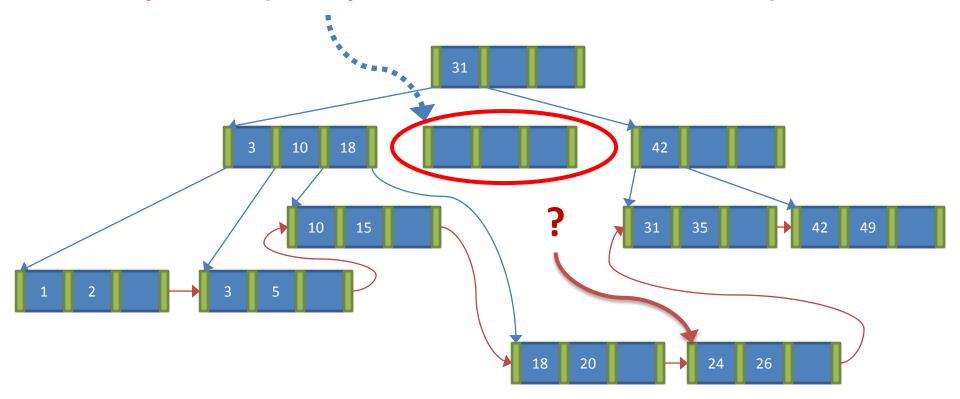
Step 2. Update the parent (Parent node is full!)
As we cannot have 5 pointers stored in a non-leaf node, we need to split this non-leaf node (Recursively).





Splitting non-leaf node (Recursive)

Step 1. We first create a new node to accommodate the new **pointers** (the 5 pointers, one for each leaf node).

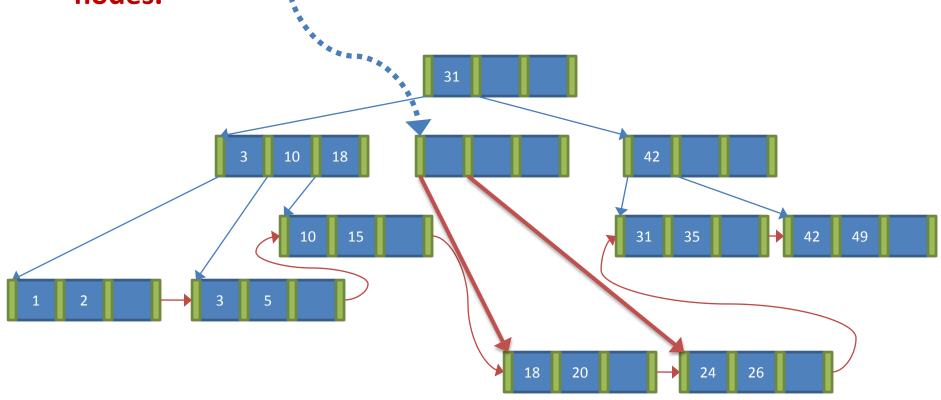




Splitting non-leaf node

Step 2. We distribute the pointers among the two

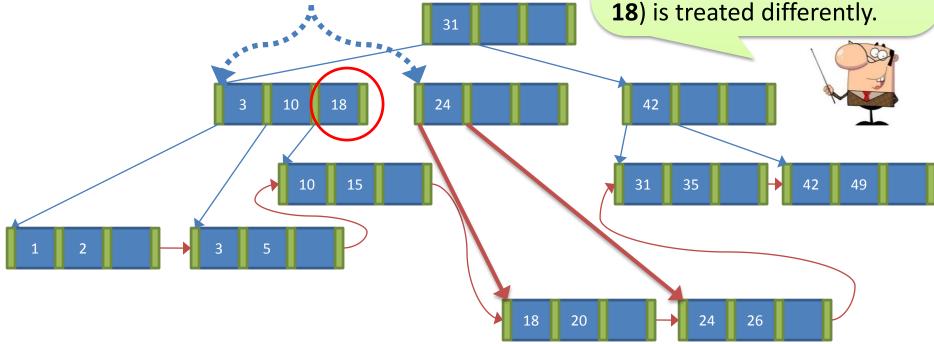






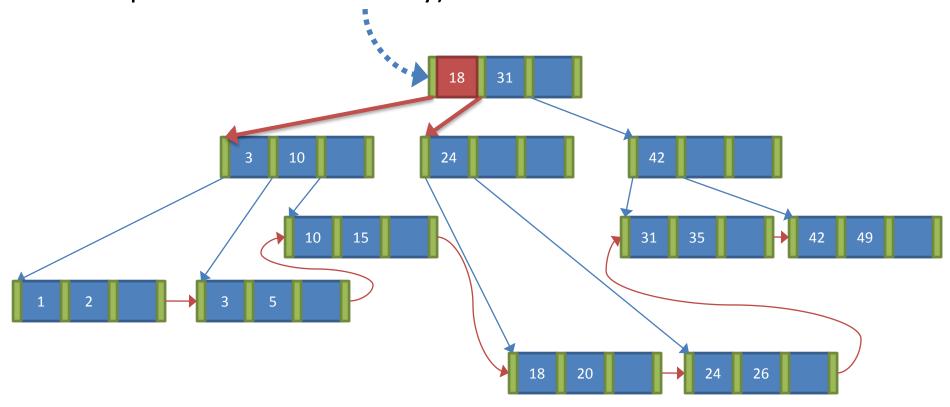
Step 3. Then consider the keys that are required in each slot among the two nodes.

Tricky part! Note that the search key that lies between the pointers that stay on the left, and the pointers that move to the right node (i.e. 18) is treated differently



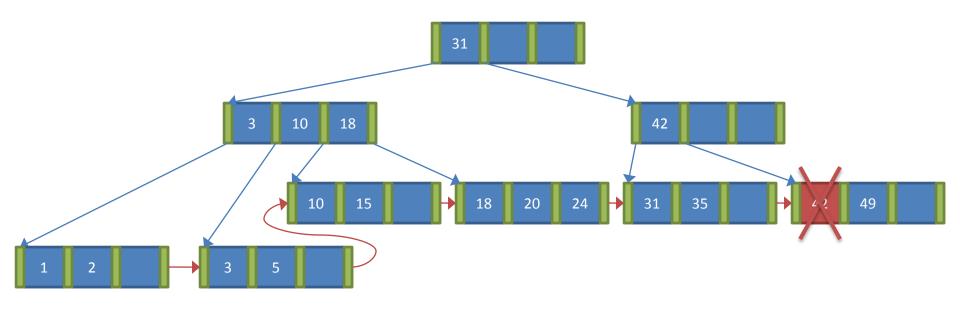


"18" is moved to the parent node to separate the searchkeys among the two nodes (if the parent node is full, split the parent node recursively)



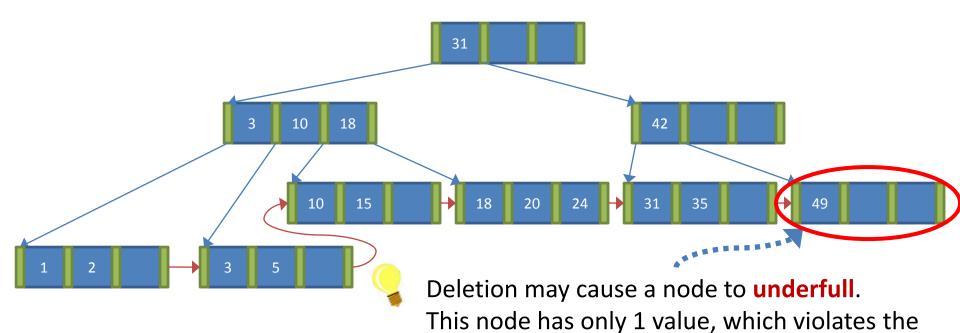
Deletion

- Find the record to be deleted.
- Remove it from the file and from the leaf node (if present)
- If the leaf node has too few entries due to the removal:
 - Try to MERGE the node with its sibling node.
 - Try to REDISTRIBUTE the entries if MERGING fails.



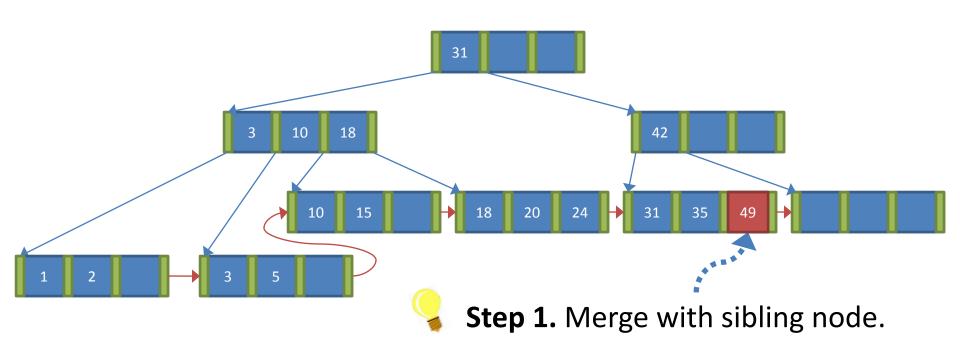


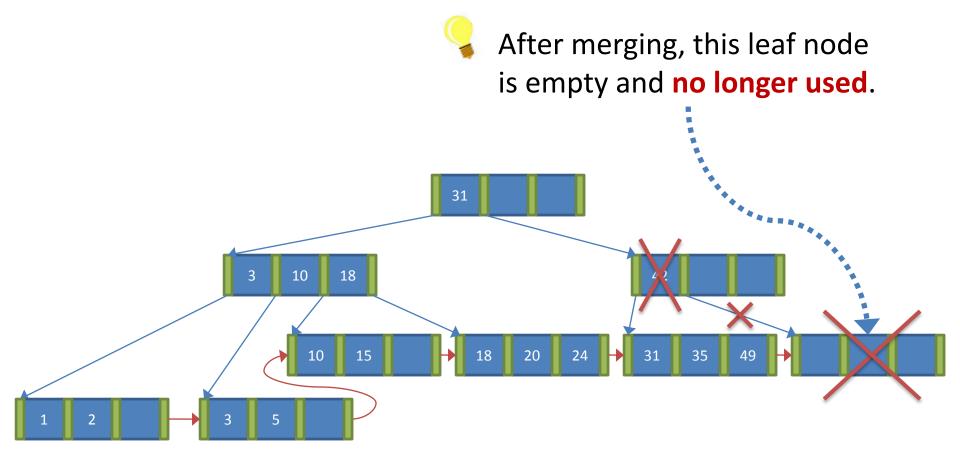
Let's try to remove key "42" in the above B+-tree.

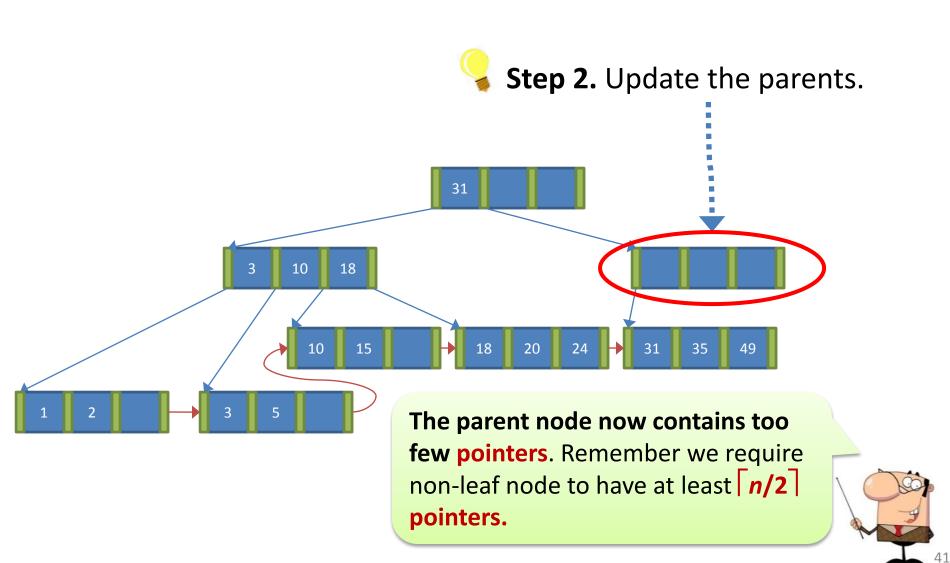


requirement that each leaf node must contain

at least $\lceil (n-1)/2 \rceil$ values (i.e., 2 in this case).

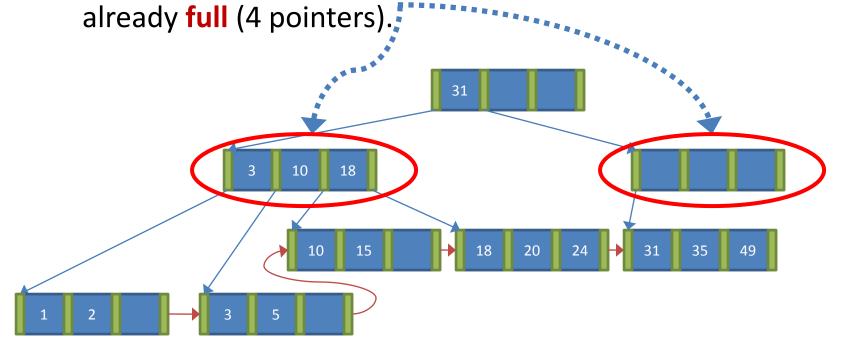






Recursively, we try to **MERGE** these 2 nodes.

However, the two nodes cannot be merged as the left node is



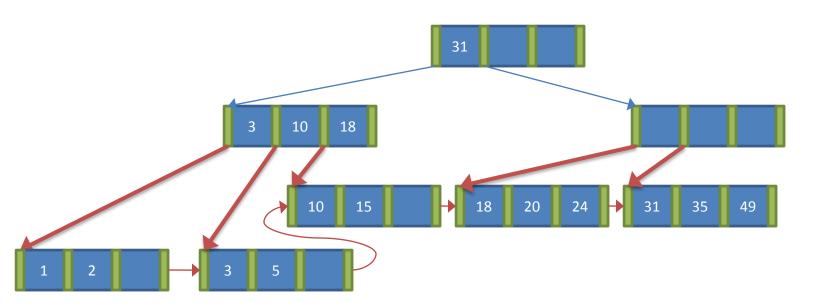
When MERGE fails, do REDISTRIBUTION!

2. Redistribution

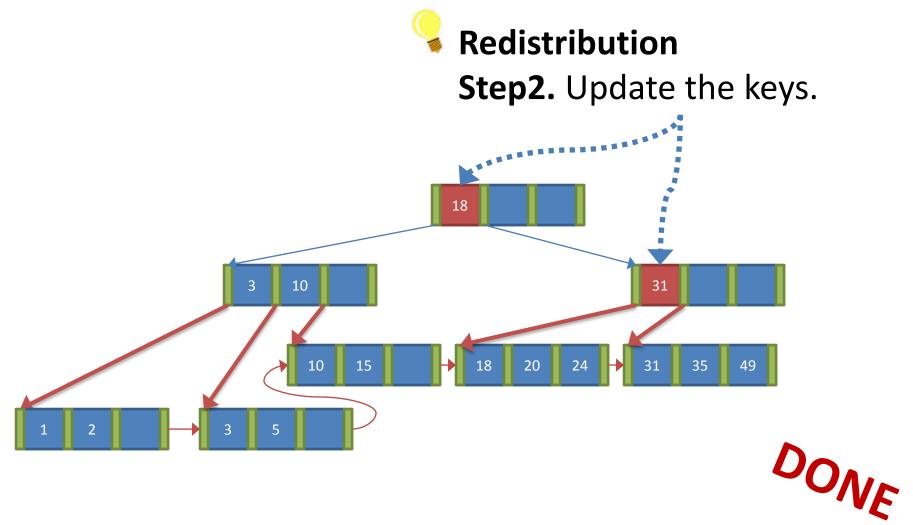


Redistribution

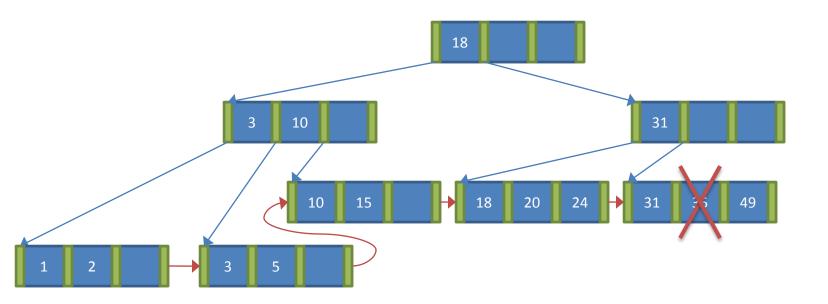
Step1. Redistribute the pointers.



2. Redistribution

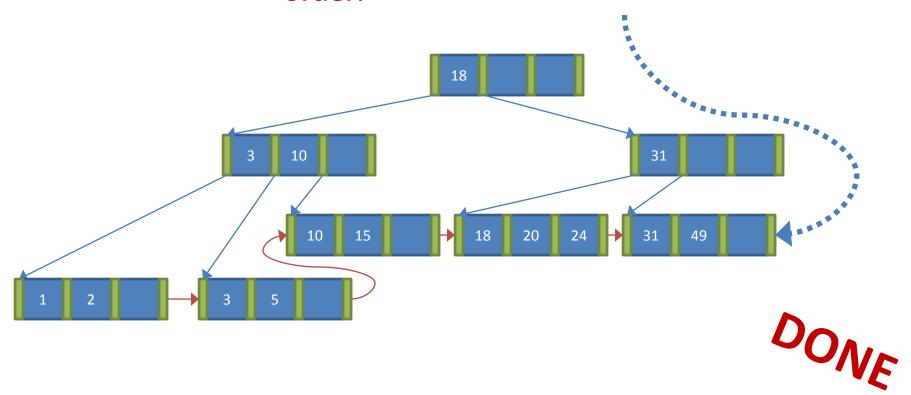


Delete 35



Delete 35

After deletion, this node contains 2 values (VALID). Remember the keys in a node should be in sorted order.

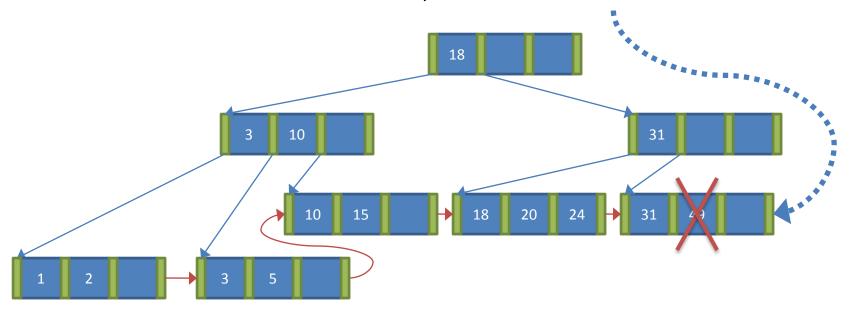


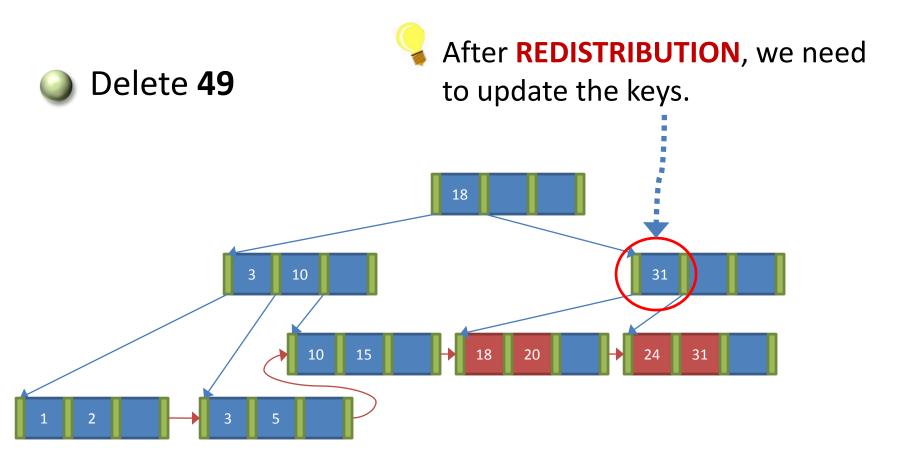
Delete 49



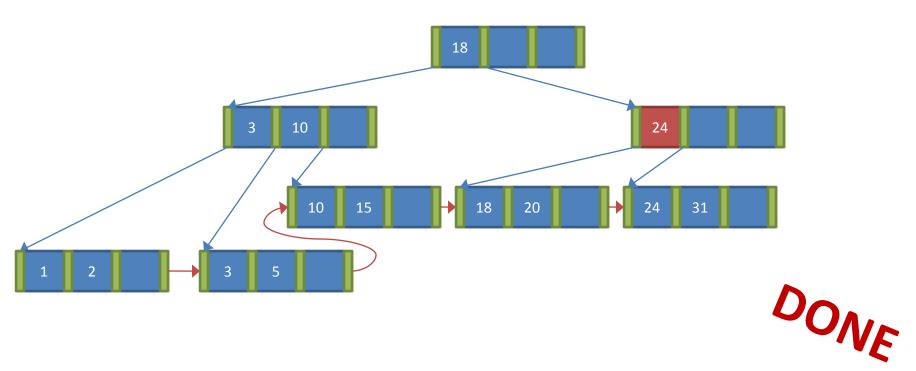
Deletion of "49" causes this leaf node to contain only one value, which is underfull.

First, try **MERGE** with its sibling node, but the sibling node is full, so we need to do REDISTRIBUTION.





Delete 49

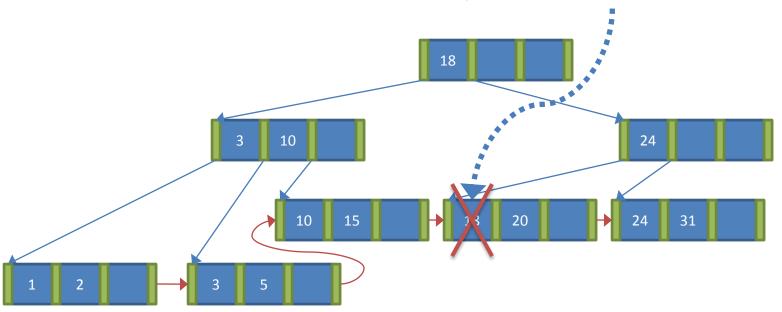




Delete **18**

Deletion of "18" causes this leaf node to contain only one value, which is underfull.

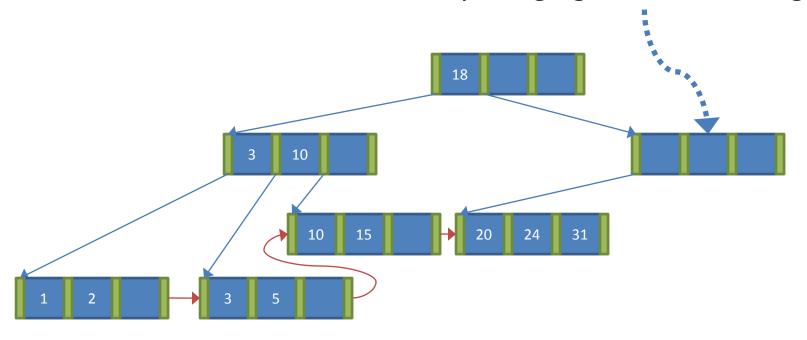
First, try merge with its sibling node, which sibling should be merged?

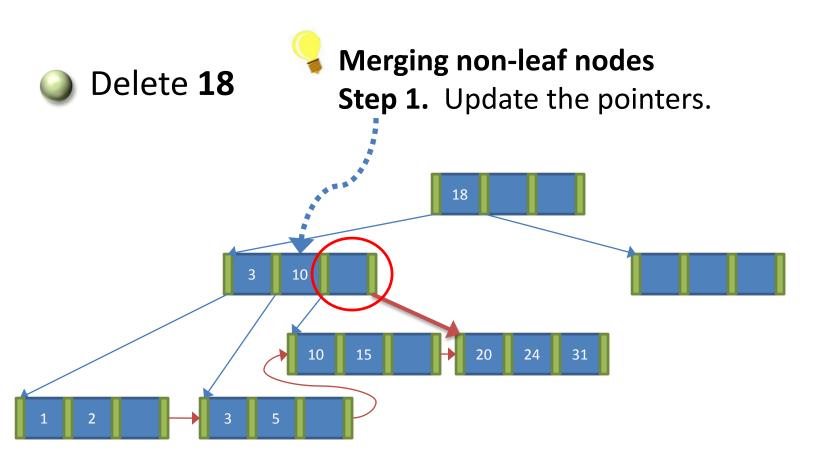


After merging, this leaf node Delete 18 is empty and no longer used. 18 10

Oelete 18

Now this node has only one pointer, which is **underfull (1 pointer only)**. We try merging it with its sibling.





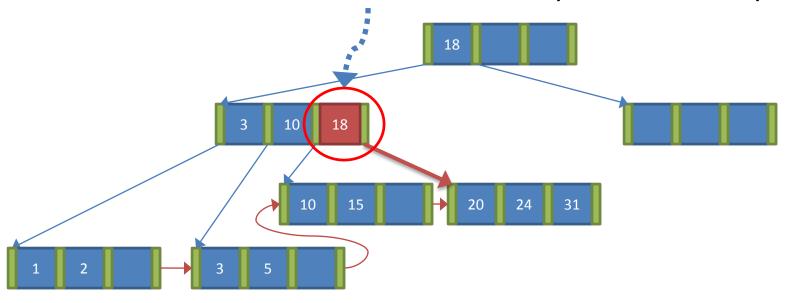




Merging non-leaf nodes

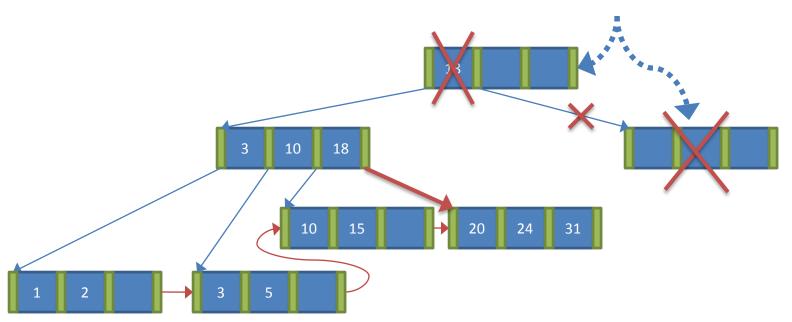
Step 2. Update the keys.

(It is "18" as originally it is the key "18" in the root node that separate the two pointers.)

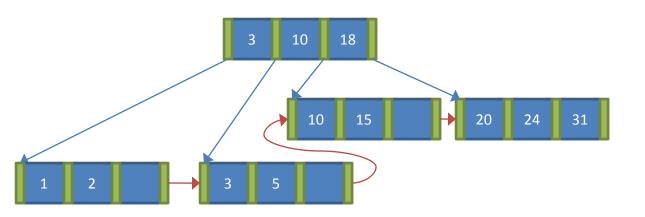


Delete 18

Note that since we merged the nonleaf node, some pointers and parent entries can be removed.



Delete 18





Section 3

Hashing

Hashing

- A kind of un-ordered indices.
 - No order within the index entries.
- The entries are divided into buckets.
- A bucket is a unit of storage containing one or more records (a bucket is typically a disk block).
- In a hash file organization, we obtain the bucket of a record directly from its search-key value using a hash function.

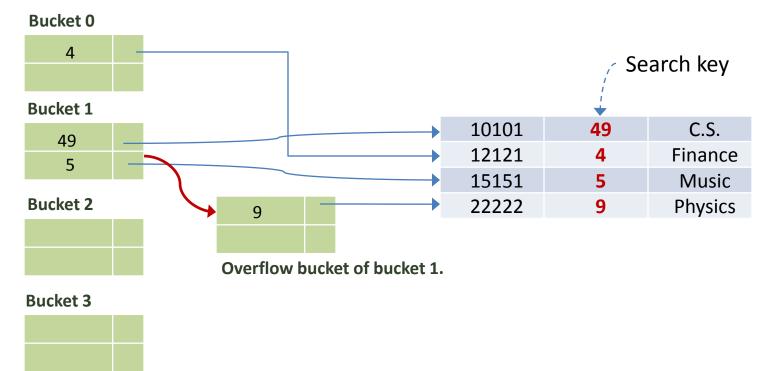
Hashing

- A hash function h() is a function that map the set of all search-key values to the set of all bucket addresses.
 - E.g., h("Peter") = i: Peter's record is mapped to bucket i.
- Records with different search-key values may be mapped to the same bucket (Collision).
 - I.e., the entire bucket has to be searched sequentially to locate a record

Static hashing

Settings

- Hash function : K mod 4.
- Each bucket can hold 2 entries.



Static hashing

Problems

- Database grows with time. If the initial number of buckets is too small, performance will degrade due to too much overflow.
- If we anticipate the database size and allocate space accordingly, a significant amount of space is wasted initially.

Solution

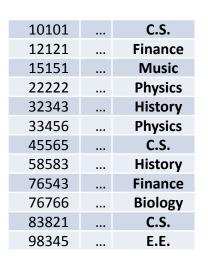
Dynamic hashing (e.g., extendable hashing)

| Department | h(department) – binary representation |
|------------|---|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| C.S. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| E.E. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |

| C.S. | ••• | 10101 |
|---------|-----|-------|
| Finance | ••• | 12121 |
| Music | ••• | 15151 |
| Physics | | 22222 |
| History | | 32343 |
| Physics | | 33456 |
| C.S. | | 45565 |
| History | ••• | 58583 |
| Finance | | 76543 |
| Biology | | 76766 |
| C.S. | | 83821 |
| E.E. | | 98345 |
| | | |

| Department | h(department) – binary representation |
|------------|---|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| C.S. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| E.E. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |

Hash prefix: The first **0** bits are used in hashing.



Assumption

Each bucket can store 2 entries.

All records in this bucket share the same **0** bit(s) in their hash prefix.

Bucket 1

Bucket address table

| Department | h(department) – binary representation |
|------------|---|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| C.S. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| E.E. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |

Assumption

Each bucket can store 2 entries.

All records in this bucket share the same **0** bit(s) in their hash prefix.

10101 ... C.s.

Bucket 1

Hash prefix: The first **0** bits are used in hashing.

| Insert | | | |
|--------|---------|-----|-------|
| | C.S. | | 10101 |
| | Finance | ••• | 12121 |
| | Music | ••• | 15151 |
| | Physics | | 22222 |
| | History | ••• | 32343 |
| | Physics | | 33456 |
| | C.S. | | 45565 |
| | History | | 58583 |
| | Finance | | 76543 |
| | Biology | | 76766 |
| | C.S. | | 83821 |
| | E.E. | | 98345 |

Bucket address table



Since currently hash prefix = 0, which means the first 0 bits are used in hashing (no bits), so simply store the record in the only bucket.

| Department | h(department) – binary representation |
|------------|---|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| C.S. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| E.E. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |

Insert

Hash prefix: The first **0** bits are used in hashing.

C.S.

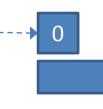
Finance

Music Physics

History

C.S.

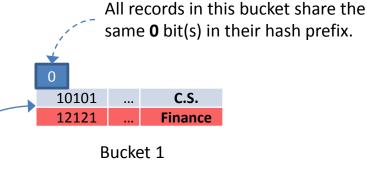
E.E.



Bucket address table

Assumption

Each bucket can store 2 entries.



| 1113601 9 | ••• | 32373 |
|-----------|-----|-------|
| Physics | ••• | 33456 |
| C.S. | ••• | 45565 |
| History | ••• | 58583 |
| Finance | ••• | 76543 |
| Biology | ••• | 76766 |
| | | |

10101

12121

15151

22222

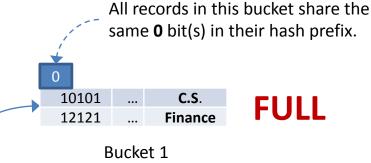
83821

98345

| Department | h(department) – binary representation |
|------------|---|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| C.S. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| E.E. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |

Insert

Hash prefix: The first **0** bits are used in hashing.



Bucket address table

| C.S. | ••• | 10101 |
|---------|-----|-------|
| Finance | ••• | 12121 |
| Music | ••• | 15151 |
| Physics | | 22222 |
| History | | 32343 |
| Physics | | 33456 |
| C.S. | | 45565 |
| History | | 58583 |
| Finance | | 76543 |
| Biology | | 76766 |
| C.S. | | 83821 |
| E.E. | | 98345 |

In extendable hashing, when a bucket is full, we need to add one more bucket. There are two options:

- 1. Add a bucket and increase the number of bits that we use from the hash value.
- **2.** Add as overflow bucket (When 1. cannot solve the problem).



| h(department) – binary representation |
|---|
| 0010 1101 1111 1011 0010 1100 0011 0000 |
| 1111 0001 0010 0100 1001 0011 0110 1101 |
| 0100 0011 1010 1100 1100 0110 1101 1111 |
| 1010 0011 1010 0000 1100 0110 1001 1111 |
| 1100 0111 1110 1101 1011 1111 0011 1010 |
| 0011 0101 1010 0110 1100 1001 1110 1011 |
| 1001 1000 0011 1111 1001 1100 0000 0001 |
| |

Hash prefix: The first **1** bit is used in hashing.

10101 C.S. 12121 **Finance** ▲ Insert Music 15151 22222 **Physics** 32343 History 33456 **Physics** 45565 C.S. 58583 History 76543 **Finance** 76766 **Biology** 83821 C.S. 98345 E.E.

0 1 Bucket address table All records in this bucket share the same 1 bit(s) in their hash prefix.

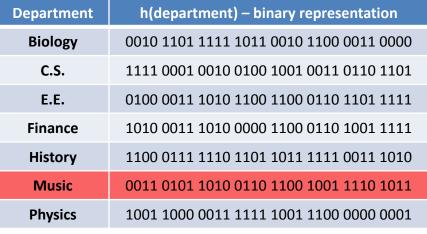
C.S.

Finance

10101

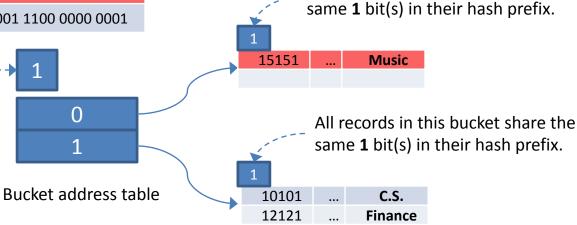
12121

All records in this bucket share the



Hash prefix: The first **1** bit is used in hashing.

10101 C.S. 12121 **Finance** ▲ Insert 15151 Music 22222 **Physics** 32343 History 33456 **Physics** 45565 C.S. 58583 History 76543 **Finance** 76766 **Biology** 83821 C.S. E.E. 98345

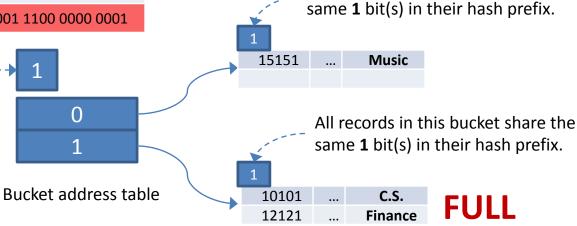


All records in this bucket share the

| h(department) – binary representation |
|---|
| 0010 1101 1111 1011 0010 1100 0011 0000 |
| 1111 0001 0010 0100 1001 0011 0110 1101 |
| 0100 0011 1010 1100 1100 0110 1101 1111 |
| 1010 0011 1010 0000 1100 0110 1001 1111 |
| 1100 0111 1110 1101 1011 1111 0011 1010 |
| 0011 0101 1010 0110 1100 1001 1110 1011 |
| 1001 1000 0011 1111 1001 1100 0000 0001 |
| |

Hash prefix: The first **1** bit is used in hashing.

10101 C.S. 12121 **Finance** 15151 Music Insert 22222 **Physics** 32343 History 33456 **Physics** 45565 C.S. 58583 History 76543 **Finance** 76766 **Biology** 83821 C.S. E.E. 98345

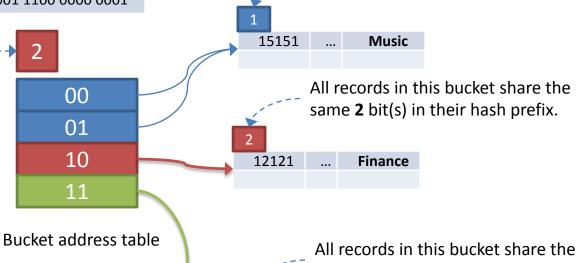


All records in this bucket share the

| Department | h(department) – binary representation |
|------------|---|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| C.S. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| E.E. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |

Hash prefix: The first **2** bits are used in hashing.

10101 C.S. 12121 **Finance** Music 15151 Insert 22222 **Physics** 32343 History 33456 **Physics** 45565 C.S. 58583 History 76543 **Finance** 76766 **Biology** 83821 C.S. E.E. 98345



10101

All records in this bucket share the same **1** bit(s) in their hash prefix.

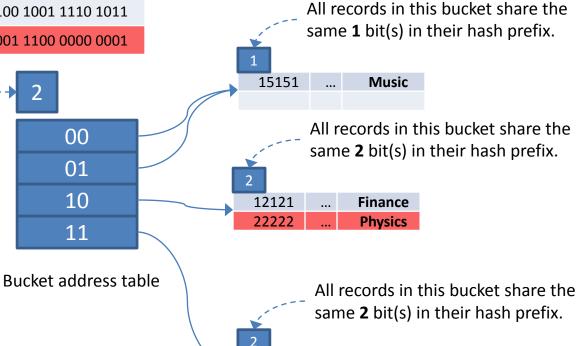
same 2 bit(s) in their hash prefix.

C.S.

| Department | h(department) – binary representation |
|------------|---|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| C.S. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| E.E. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |

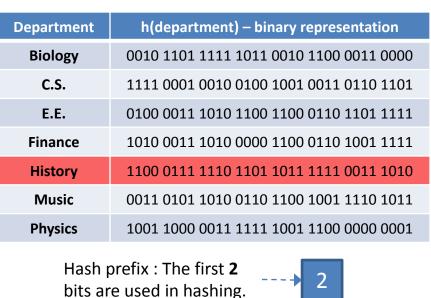
Hash prefix: The first **2** bits are used in hashing.

10101 C.S. 12121 **Finance** 15151 Music Insert 22222 **Physics** 32343 History 33456 **Physics** 45565 C.S. 58583 History 76543 **Finance** 76766 **Biology** 83821 C.S. E.E. 98345

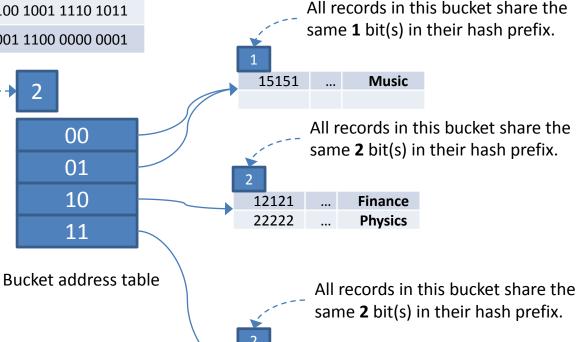


10101

C.S.



C.S. 10101 12121 Finance 15151 Music 22222 **Physics** Insert 32343 History 33456 **Physics** 45565 C.S. 58583 History 76543 **Finance** 76766 **Biology** 83821 C.S. E.E. 98345



10101

32343

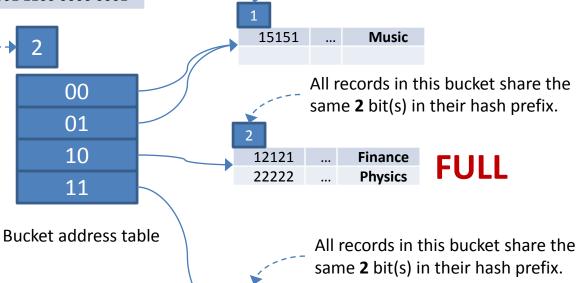
C.S.

History

| Department | h(department) – binary representation |
|------------|---|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| C.S. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| E.E. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |

Hash prefix: The first **2** bits are used in hashing.

10101 C.S. 12121 **Finance** 15151 Music 22222 **Physics** 32343 History Insert 33456 **Physics** 45565 C.S. 58583 History 76543 **Finance** 76766 **Biology** 83821 C.S. E.E. 98345



10101

32343

C.S.

History

73

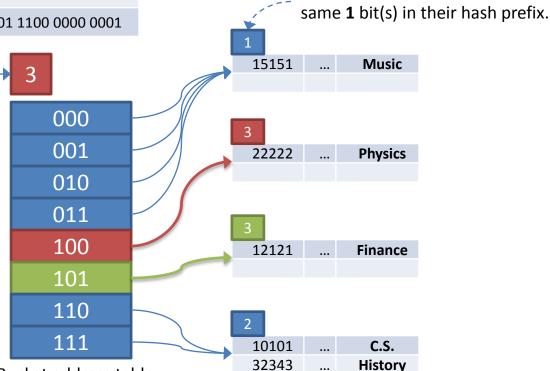
All records in this bucket share the same **1** bit(s) in their hash prefix.

Bucket address table

| Department | h(department) – binary representation |
|------------|---|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| C.S. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| E.E. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |
| | |

Hash prefix: The first **3** bits are used in hashing.

10101 C.S. 12121 **Finance** 15151 Music 22222 **Physics** 32343 History Insert 33456 **Physics** C.S. 45565 58583 History 76543 **Finance** 76766 **Biology** 83821 C.S. 98345 E.E.

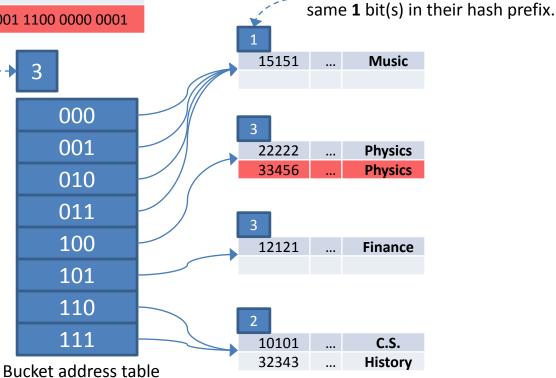


All records in this bucket share the

| Department | h(department) – binary representation |
|---------------------------|---|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| C.S. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| E.E. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |
| Hash prefix : The first 3 | |

bits are used in hashing.

| | C.S. | ••• | 10101 |
|-----------------|----------------|-----|-------|
| | Finance | ••• | 12121 |
| | Music | | 15151 |
| | Physics | ••• | 22222 |
| ⊿ Insert | History | | 32343 |
| III.SCI C | Physics | | 33456 |
| | C.S. | | 45565 |
| | History | ••• | 58583 |
| | Finance | ••• | 76543 |
| | Biology | | 76766 |
| | C.S. | ••• | 83821 |
| | E.E. | | 98345 |

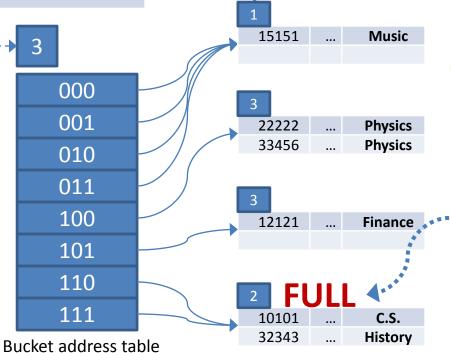


All records in this bucket share the

| Department | h(department) – binary representation |
|------------|---|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| C.S. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| E.E. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |

Hash prefix: The first **3** bits are used in hashing.

| | C.S. | ••• | 10101 |
|-----------------|---------|-----|-------|
| | Finance | | 12121 |
| | Music | ••• | 15151 |
| | Physics | ••• | 22222 |
| | History | ••• | 32343 |
| ₄ Insert | Physics | ••• | 33456 |
| | C.S. | ••• | 45565 |
| | History | | 58583 |
| | Finance | ••• | 76543 |
| | Biology | ••• | 76766 |
| | | | |
| | C.S. | | 83821 |



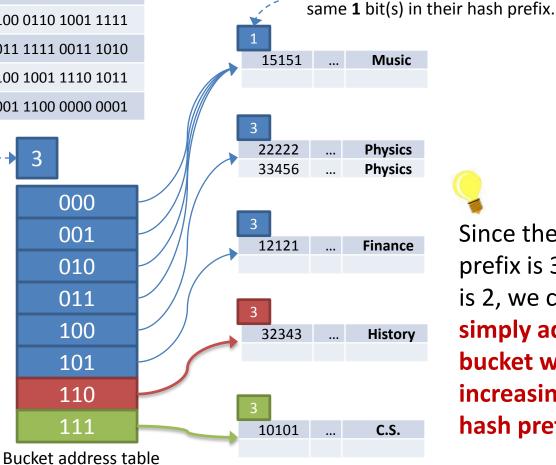
Since the hash prefix is 3 and 2 is 2, we can simply add a bucket without increasing the hash prefix.

All records in this bucket share the same **1** bit(s) in their hash prefix.

| Department | h(department) – binary representation |
|------------|---|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| C.S. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| E.E. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |

Hash prefix: The first 3 bits are used in hashing.

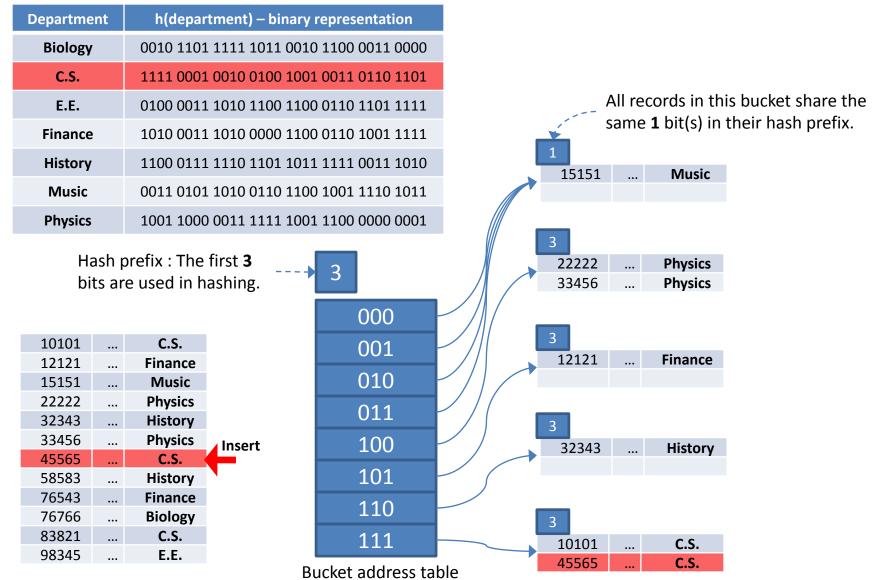
| | C.S. | | 10101 |
|-----------------|---------|-----|-------|
| | Finance | ••• | 12121 |
| | Music | ••• | 15151 |
| | Physics | ••• | 22222 |
| | History | ••• | 32343 |
| ₄ Insert | Physics | ••• | 33456 |
| | C.S. | | 45565 |
| | History | | 58583 |
| | Finance | ••• | 76543 |
| | Biology | ••• | 76766 |
| | C.S. | ••• | 83821 |
| | F.F. | | 98345 |

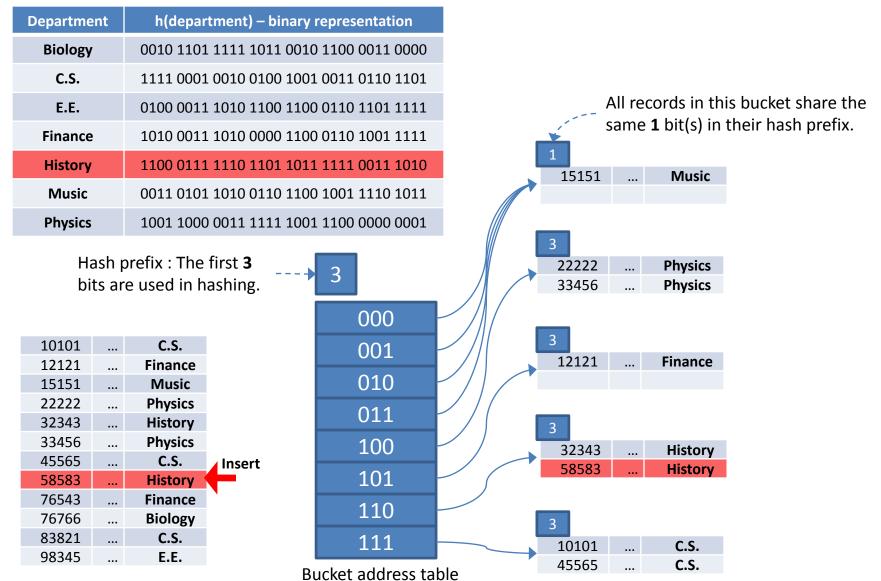


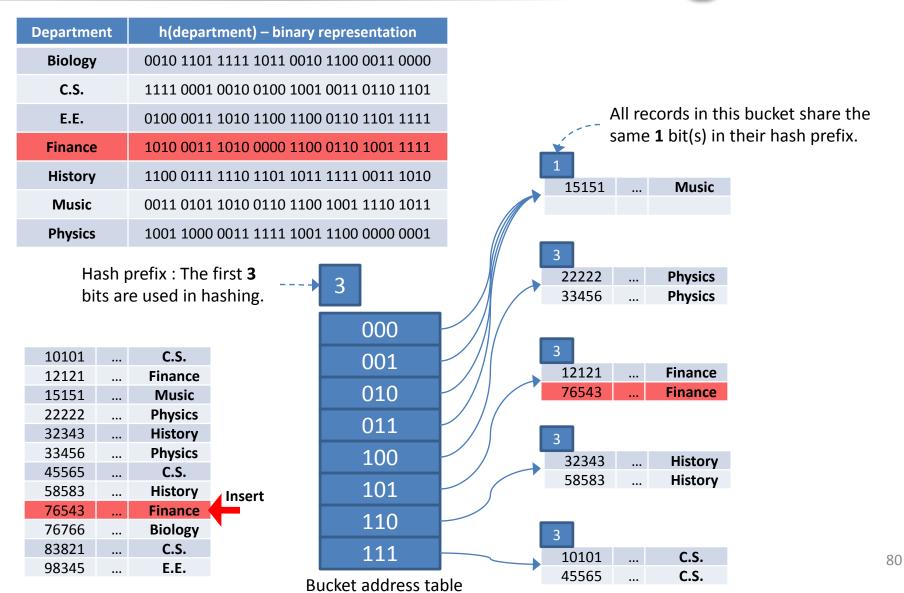


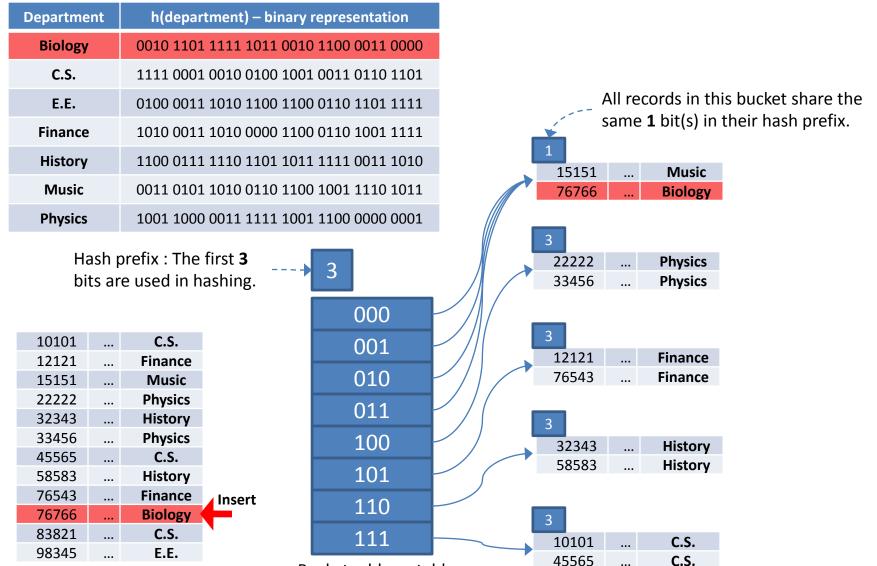
All records in this bucket share the

Since the hash prefix is 3 and 2 is 2, we can simply add a **bucket without** increasing the hash prefix.









Bucket address table

000

001

010

011

100

101

110

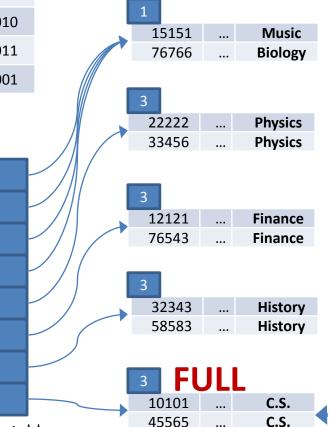
111

Bucket address table

| Department | h(department) – binary representation |
|------------|---|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| C.S. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| E.E. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |

Hash prefix: The first **3** bits are used in hashing.

| | C.S. | ••• | 10101 |
|-----------------|---------|-----|-------|
| | Finance | ••• | 12121 |
| | Music | ••• | 15151 |
| | Physics | | 22222 |
| | History | ••• | 32343 |
| | Physics | ••• | 33456 |
| | C.S. | ••• | 45565 |
| | History | | 58583 |
| | Finance | ••• | 76543 |
| ⊿ Insert | Biology | | 76766 |
| | C.S. | | 83821 |
| • | E.E. | | 98345 |



Note that all entries in this bucket are of the same hash value "C.S.", increasing hash prefix will not solve the collision problem!

Chaining is used.

82

All records in this bucket share the same **1** bit(s) in their hash prefix.

000

001

010

011

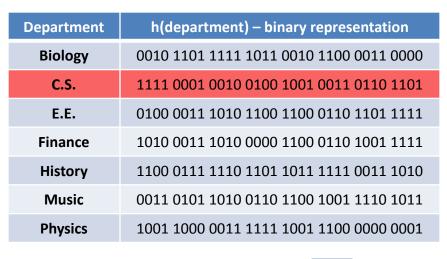
100

101

110

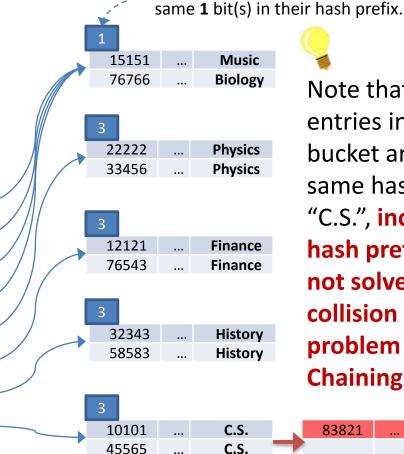
111

Bucket address table



Hash prefix: The first 3 bits are used in hashing.

| | C.S. | | 10101 |
|-----------------|---------|-----|-------|
| | Finance | | 12121 |
| | Music | | 15151 |
| | Physics | | 22222 |
| | History | | 32343 |
| | Physics | | 33456 |
| | C.S. | | 45565 |
| | History | | 58583 |
| | Finance | | 76543 |
| ⊿ Insert | Biology | ••• | 76766 |
| - inscre | C.S. | | 83821 |
| • | E.E. | | 98345 |

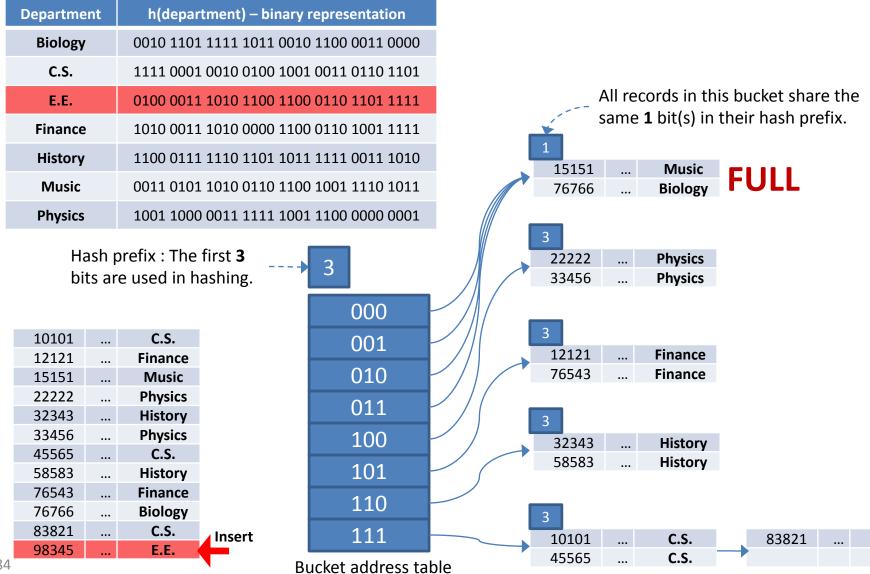


Note that all entries in this bucket are of the same hash value "C.S.", increasing hash prefix will not solve the collision problem! Chaining is used.

C.S.

83821

All records in this bucket share the

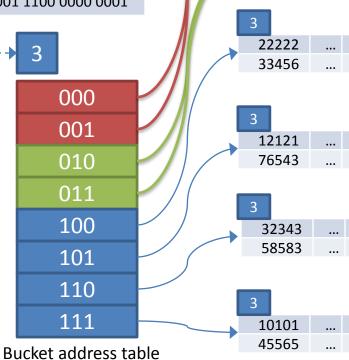


C.S.

| Department | h(department) – binary representation |
|------------|---|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| C.S. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| E.E. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |

Hash prefix: The first **3** bits are used in hashing.

| C.S. | 10101 |
|---------|-----------|
| Finance | 12121 |
| Music | 15151 |
| Physics | 22222 |
| History | 32343 |
| Physics | 33456 |
| C.S. | 45565 |
| History | 58583 |
| Finance | 76543 |
| Biology | 76766 |
| C.S. | 83821 |
| E.E. | 98345 |
| | |



15151

76766

All records in this bucket share the same **2** bit(s) in their hash prefix.

Music

Biology

Physics

Physics

Finance

Finance

History

History

C.S.

C.S.

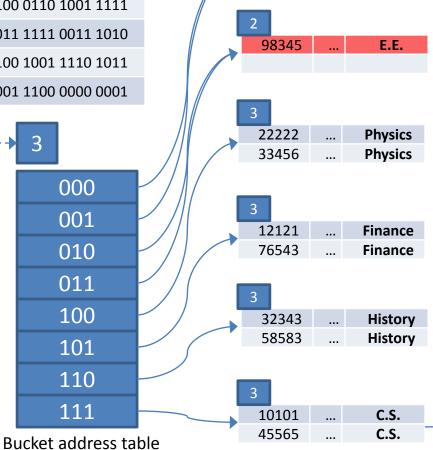
83821

C.S.

| Department | h(department) – binary representation |
|------------|---|
| Biology | 0010 1101 1111 1011 0010 1100 0011 0000 |
| C.S. | 1111 0001 0010 0100 1001 0011 0110 1101 |
| E.E. | 0100 0011 1010 1100 1100 0110 1101 1111 |
| Finance | 1010 0011 1010 0000 1100 0110 1001 1111 |
| History | 1100 0111 1110 1101 1011 1111 0011 1010 |
| Music | 0011 0101 1010 0110 1100 1001 1110 1011 |
| Physics | 1001 1000 0011 1111 1001 1100 0000 0001 |

Hash prefix: The first **3** bits are used in hashing.

| C.S. | ••• | 10101 |
|---------|-----|-------|
| Finance | | 12121 |
| Music | | 15151 |
| Physics | | 22222 |
| History | | 32343 |
| Physics | | 33456 |
| C.S. | | 45565 |
| History | | 58583 |
| Finance | | 76543 |
| Biology | | 76766 |
| C.S. | | 83821 |
| E.E. | | 98345 |
| | | |



15151

76766

All records in this bucket share the same **2** bit(s) in their hash prefix.

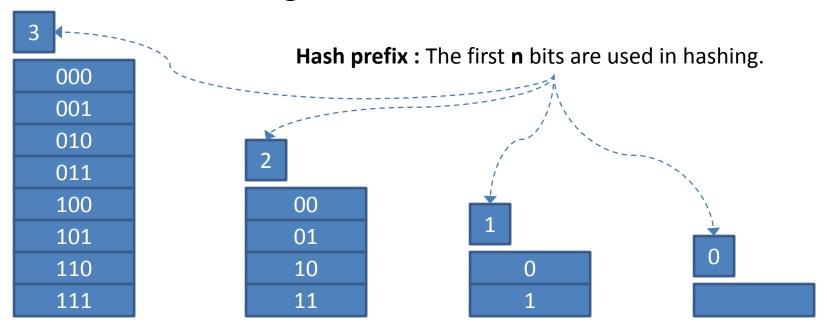
C.S.

83821

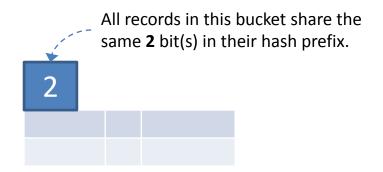
Music

Biology

- The bucket address table has a hash prefix value n.
- The first-*n* bits of the hash values are used to locate the address table entry for the pointer pointing to the bucket containing records with that hash value.



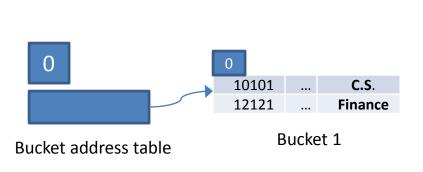
- Each bucket j contains:
 - The records
 - 🍑 An integer **i**j

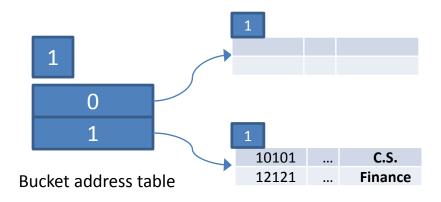


All records in the jth bucket share the same first-i_j bits in the hash function.

Splitting a bucket

- If the hash prefix = i_j
 - Only 1 pointer in address table is pointing to bucket j.
 - We increment hash prefix by 1, and double the size of the bucket address table.
 - Reinsert the entries in the original bucket.





Splitting a bucket

- If the hash prefix $> i_j$ (i.e., the indicator in the bucket address table is larger than the indicator in bucket)
 - \bigcirc Create a new bucket $oldsymbol{z}$, and set $oldsymbol{i_j}$ and $oldsymbol{i_z}$ to old $oldsymbol{i_j}$ + 1.
 - There must be a number of address table entries pointing to bucket j originally, change the lower half of such entries so that they point to bucket z.
 - Remove and re-insert each record in bucket j.

Let's look at the running example: the insertion of the entry "45565,...,C.S." for an illustration.



Hash indexes v.s. B+-tree

- Hash indexes have different characteristics than B+-tree
 - They are used only for equality comparisons that use the = operator (but they are very fast).
 - They are not used for comparison operators such as < that find a range of values.
 - The optimizer cannot use a hash index to speed up ORDER BY operations.

Section 4

Index & SQL

Defining index in SQL

To create an index:

```
CREATE INDEX <index-name> ON
  <relation-name> ( <attribute-list> )
  [index_type]
```

- Optional [index_type]: USING {BTREE | HASH}
- Use CREATE UNIQUE INDEX to indirectly specify and enforce the condition that the search-key is a superkey.
- To remove an index

DROP INDEX <index-name>

Chapter 7.

END

CSIS0278 / COMP3278 Introduction to Database Management Systems

Department of Computer Science, The University of Hong Kong

Slides prepared by - **Dr. Chui Chun Kit**, http://www.cs.hku.hk/~ckchui/ for students in CSIS0278 / COMP3278 For other uses, please email : ckchui@cs.hku.hk